

# Advanced Programming

Generator functions

# READING A LARGE FILE

- ✦ Assume that we have a very large text file containing billions of lines
- ✦ Process this file line by line requires to
  - ✦ Load the entire file into memory
    - ✦ Inefficient and may not fit into memory
  - ✦ Tokenize each line and store them in an ADT (lists or dictionary)
    - ✦ Inefficient and may not fit into memory

```
import re

def process_file(input_file, output_file):
    with open(input_file, 'r') as f_input:
        lines = f_input.readlines()
        processed_lines = [tokenize_line(line)
                           for line in lines]
    with open(output_file, 'w') as f_output:
        for line in processed_lines:
            f_output.write(','.join(line) + '\n')

def tokenize_line(line):
    tokens = []
    if len(line) > 2:
        tokens = re.findall(r'\b\w+\b', line)
    return tokens
```

# USING A CALLBACK FUNCTION

---

- ✦ Function passed as an argument to another function
- ✦ Executed at a later time usually through an event
- ✦ They are commonly used for asynchronous programming
- ✦ Executes code concurrently and handle events efficiently
- ✦ Triggered on custom events based on conditions

```
def tokenize_file(filename, callback_fn):  
    with open(filename) as f:  
        for line in f:  
            for token in line.split():  
                callback_fn(token)  
  
def print_token(token):  
    print(f"{token}")  
  
if __name__ == "__main__":  
    tokenize_file("t2.txt", print_token)
```

Python, functions are first-class objects - they can be passed as arguments to other functions

# ITERATOR APPROACH

---

- ✦ Make tokenizer as an iterator,
  - ✦ Use iterator's `.next()` method to get the next token
  - ✦ More user-friendly
  - ✦ Burdens tokenizer with the task of maintaining state between invocations

```
import re
```

```
class Tokenizer:
```

```
    def __init__(self, input_file):  
        self.input_file = input_file  
        self.file_handle = open(input_file, 'r')  
        self.current_line = None  
        self.current_position = 0
```

```
    def tokenize_next(self):  
        if self.current_line is None or \  
        self.current_position >= len(self.current_line):  
            self.current_line = self.file_handle.readline()  
            if not self.current_line:  
                raise StopIteration  
            self.current_position = 0
```

```
        tokens = re.findall(r'\b\w+\b',  
                             self.current_line[self.current_position:])
```

```
        if not tokens:  
            # No more tokens in the current line,  
            # move to the next line  
            self.current_line = None  
            return self.tokenize_next()
```

```
        # Update the position for the next call
```

```
        token_length = len(tokens[0])
```

```
        self.current_position += token_length
```

```
        # Account for spaces between tokens
```

```
        while self.current_position < len(self.current_line) and \  
              not self.current_line[self.current_position].isalnum():  
            self.current_position += 1
```

```
        return tokens[0]
```

```
if __name__ == '__main__':  
    input_file = 't2.txt'  
    tokenizer = Tokenizer(input_file)  
    try:  
        while True:  
            token = tokenizer.tokenize_next()  
            print(token)  
    except StopIteration:  
        pass  
    tokenizer.file_handle.close()
```

# THREAD-BASED TOKENISATION

---

- ✦ Running the producer and consumer in separate threads allows both to maintain their states naturally

```
import re
import multiprocessing

def tokenize_chunk(chunk, output_queue):
    tokens = []
    for line in chunk:
        tokens.extend(tokenize_line(line))
    output_queue.put(tokens)

def main(input_file, output_file,
         num_processes=4, chunk_size=1000):
    output_queue = multiprocessing.Queue()

    with open(input_file, 'r') as f_input:
        chunks = [f_input.readlines(chunk_size)
                  for _ in range(num_processes)]

    processes = []
    for chunk in chunks:
        process = multiprocessing.Process(target=tokenize_chunk,
                                         args=(chunk, output_queue))

        process.start()
        processes.append(process)

    tokens = []
    for _ in range(num_processes):
        tokens.extend(output_queue.get())

    for process in processes:
        process.join()

    with open(output_file, 'w') as f_output:
        for token in tokens:
            f_output.write(token + '\n')
```

```
def tokenize_line(line):
    return re.findall(r'\b\w+\b', line)
```

```
if __name__ == "__main__":
    input_file = 't2.txt'
    output_file = 'output.txt'
    main(input_file, output_file)
```

# IS THERE ANY ALTERNATE SOLUTION?

---

- ✦ Develop a functionality that
  - ✦ Returns an intermediate result
  - ✦ Gets the value to its caller (consumer)
  - ✦ Maintains the local state
  - ✦ Resumes right where it stopped the computation

```
def fibonacci():  
    a, b = 0, 1  
    while True:  
        yield a  
        a, b = b, a + b  
  
if __name__ == '__main__':  
    fib = fibonacci()  
    for _ in range(10):  
        print(next(fib))
```



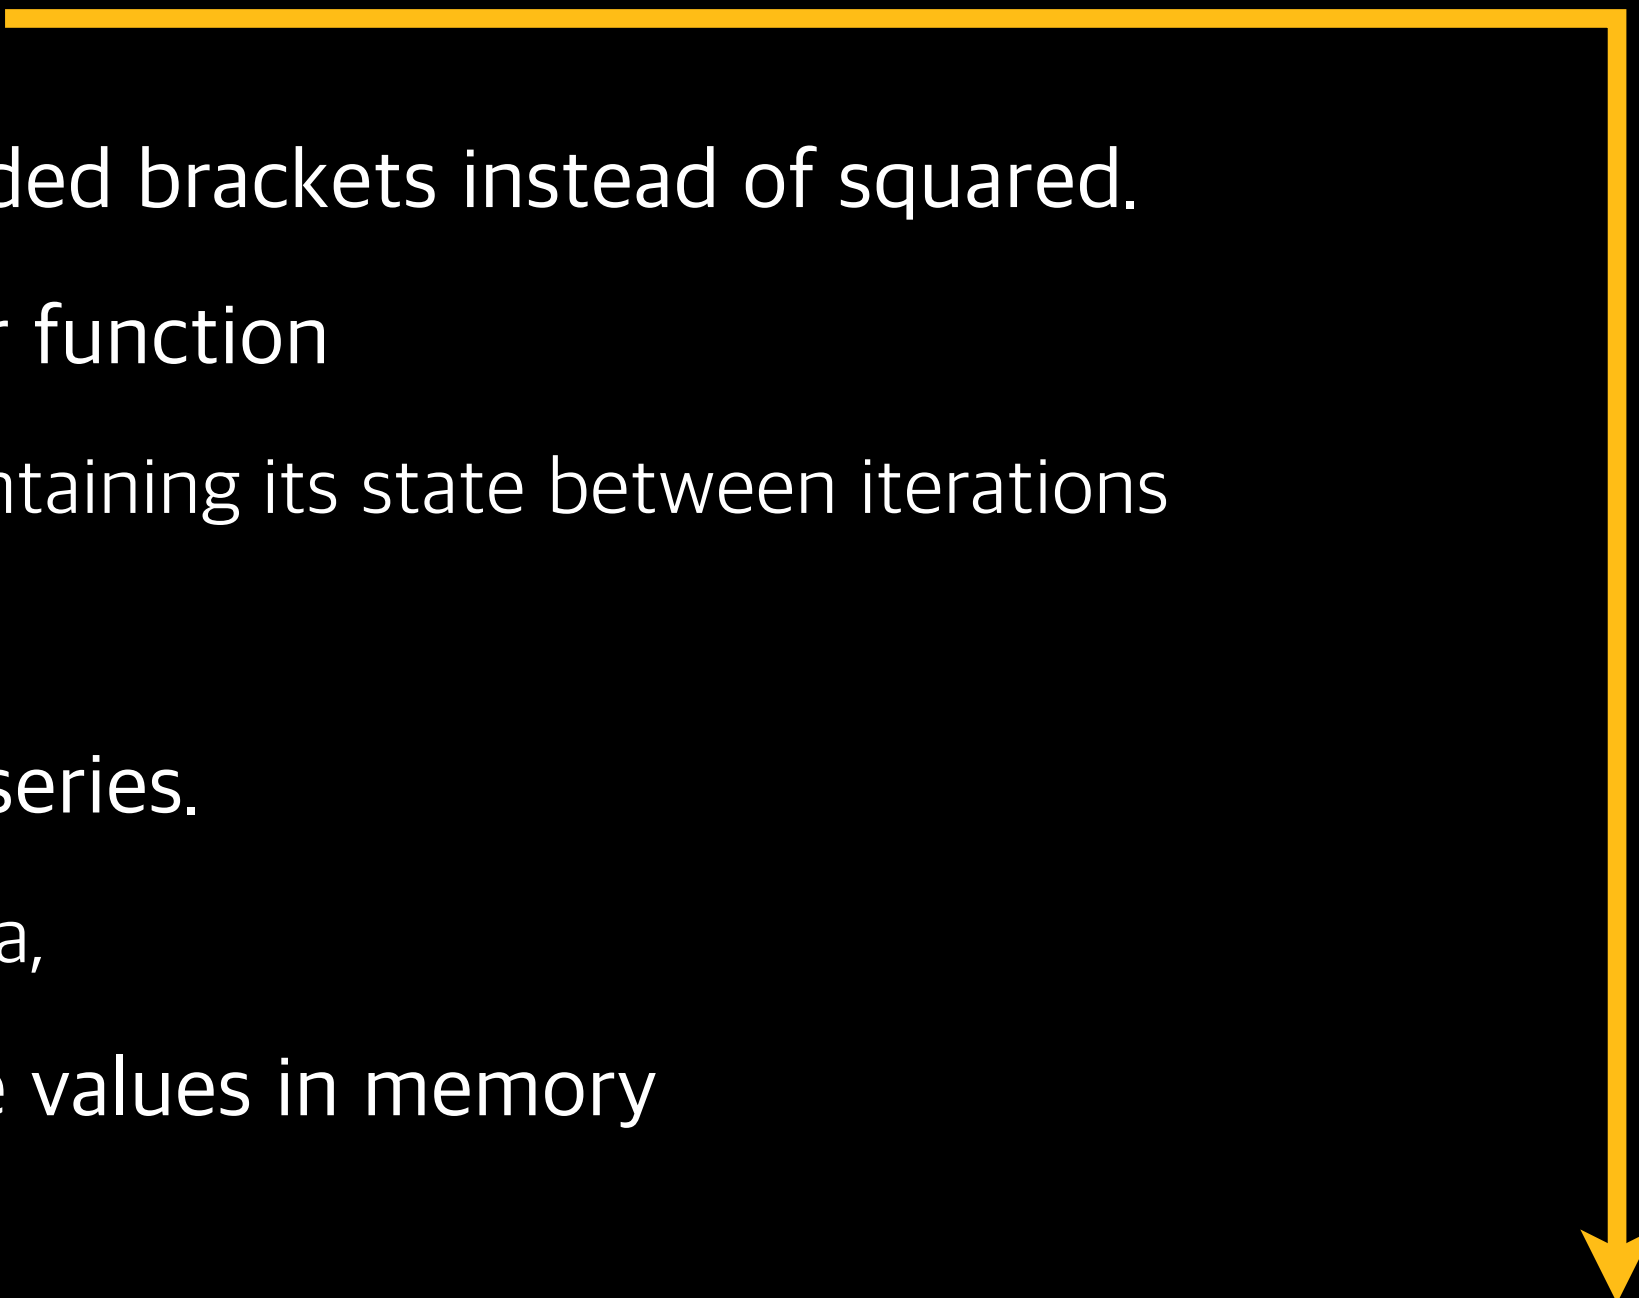
# WHAT IS A GENERATOR FUNCTION?

---

- ✦ A function that iterates over a potentially infinite sequence of values without creating the entire sequence in memory at once

# GENERATOR FUNCTION

- ✦ Uses the `yield` keyword to return an iterator that may be iterated over, one value at a time
- ✦ Generators do not store their contents in memory
  - ✦ More memory-efficient for large data sets
- ✦ Generators can be created using a generator expression
  - ✦ Similar syntax to a list comprehension but uses rounded brackets instead of squared.
  - ✦ The `yield` keyword controls the flow of a generator function
    - ✦ Allows the execution to pause and resume while maintaining its state between iterations
- ✦ Generators can be used in loops
  - ✦ Use `next()` function to retrieve the next value in the series.
- ✦ Generators can be used to represent infinite streams of data,
  - ✦ Example - Fibonacci sequence without storing all the values in memory
- ✦ Generators use lazy evaluation,



```
def squares_generator(n):  
    yield from (x ** 2 for x in range(1, n + 1))
```

```
unique_words = set(word for word in line.split())
```

# RETURN VS. YIELD

---

- ✦ Return terminates the function's execution and returns a single value
- ✦ Yield pauses execution and allows for multiple values
- ✦ Return discards the function's state after execution
- ✦ Yield preserves the state for resumption

# BENEFITS OF YIELD

---

- ✦ Memory Efficiency
  - ✦ Generator functions are memory-efficient when dealing with large sequences.
  - ✦ They only generate the next value when needed
- ✦ Lazy Evaluation
  - ✦ Values are calculated only when required, improving performance for computationally expensive sequences.
- ✦ Iterators – Generator functions create iterators

# USE CASES

---

- ✦ Infinite Sequences
  - ✦ Generating infinite sequences like prime numbers or Fibonacci numbers where you don't know the size beforehand.
- ✦ Large Datasets
  - ✦ Working with massive datasets where storing everything in memory is impractical.
- ✦ Coroutines
  - ✦ Implementing cooperative multitasking where multiple functions can yield control and resume based on events.

# TOKENIZER USING A GENERATOR

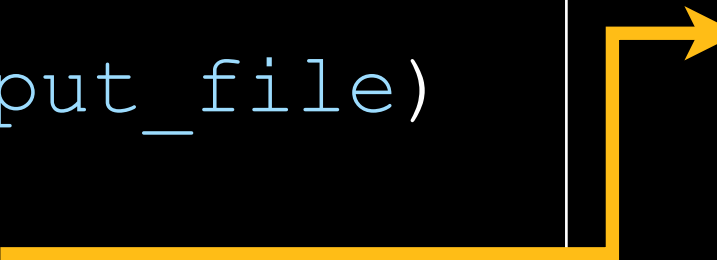
---

```
import re

def tokenizer(input_file):
    with open(input_file, 'r') as f_input:
        for line in f_input:
            for token in tokenize_line(line):
                yield token

def tokenize_line(line):
    tokens = re.findall(r'\b\w+\b', line)
    return tokens

if __name__ == '__main__':
    input_file = 't2.txt'
    token_generator = tokenizer(input_file)
    for token in token_generator:
        print(token)
```



```
try:
    while True:
        token = next(token_generator)
        print(token)
except StopIteration:
    pass
```